

HW3a (Transformers) Recitation

November 4, 2022

10-417/617 Intermediate Deep Learning

Agenda

1. HW3a Overview
2. Written portion
 - We won't really go through this in detail, but we will provide a few tips, reminders, things to watch out for
3. Transformer architecture
4. Decoding (beam search)
5. Evaluation (BLEU score)
6. Programming portion

HW3a Overview

- Released: Wednesday (November 2, 2022)
- Due: Monday (November 14, 2022)
- 80 points for 417; 100 points for 617
- Written (30+10 points):
 - 2 questions for all (15 points each)
 - 1 questions for 617 only (10 points)
- Programming (50+10 points):
 - Auto-grader (24 + 6 points)
 - Experiments / analysis (26 + 4 points)
- **Start early!**

Written portion

Q1: Vanishing/Exploding Gradients in RNNs

- Main ideas/tools: derivatives, chain rule

Q2: Deriving PyTorch's GELU

- Part 1 → change of variables (integration)
- Part 4 → make sure you remove all terms with higher order than 3 at the very first step

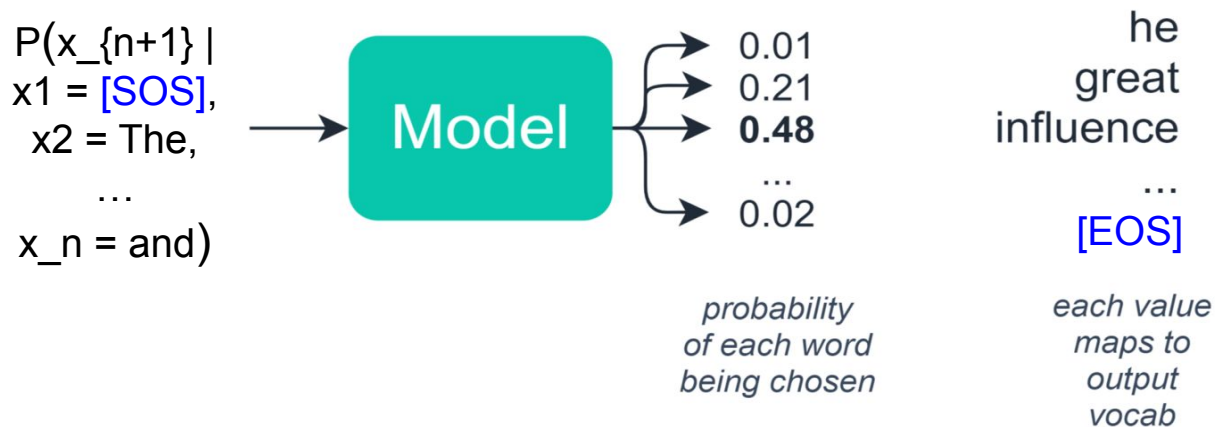
Written portion

Q3: Gradients in RBMs

- Part 2: $\beta(v)$ and $\gamma(v, h)$ are functions for you to define. Make sure you define these correctly
- Part 3: θ is just a general variable. Apply chain rule
- Part 4: Use result from part 3 (and fill in θ accordingly)

Language Modeling

A language model is a **probability distribution over a set words**, usually represented by a neural network



Note the special [SOS] and [EOS] tokens here. These are used to represent the starts and ends of sentences. In the assignment, they are index 0 and 1 in our vocabulary.

Applications of Language Modeling


Since we can predict the probability of the next word, we can use language models in open-ended generation tasks and seq2seq tasks:

- Text generation (e.g. essay generation, story generation, etc.)
- Summarization
- **Machine translation**

In this HW, we will be working on the task of French-to-English translation

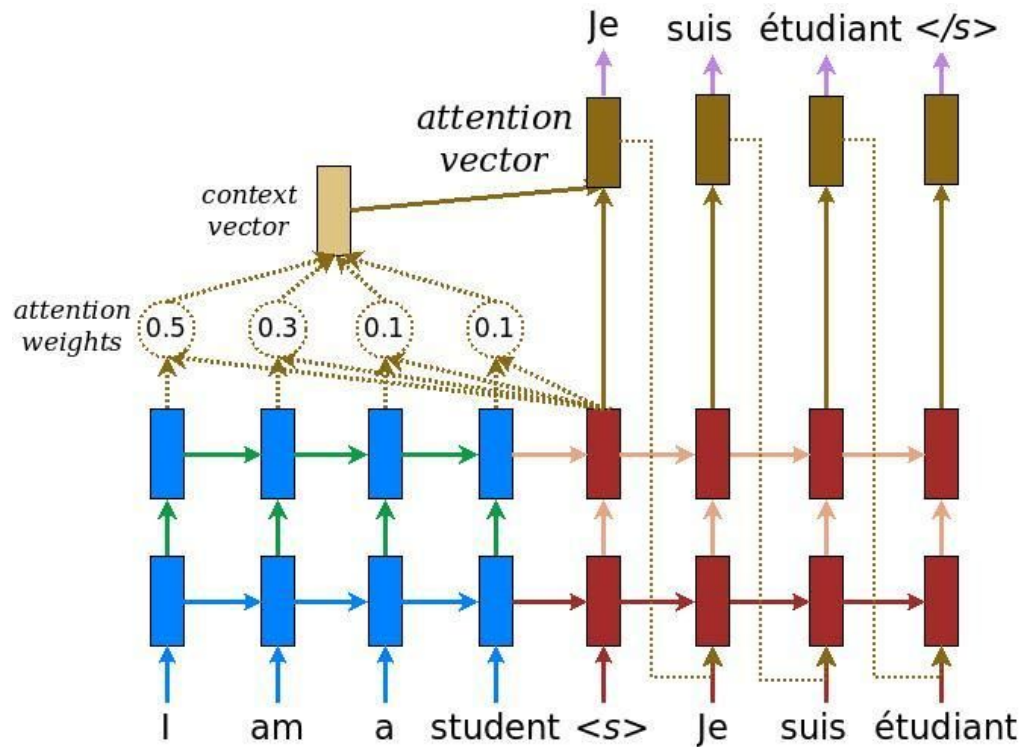
Language Modelling with RNNs

- Seq2seq translation example
- Model generates words sequentially
- One main downside is that the effect of the earlier words gets diminished over time (all the input information is encoded into just one single vector)




Attention in RNNs


- **Attention** is a key way to address this issue
 1. Take the pairwise **dot products** between the current vector and each input vector to get attention weights
 2. Multiple the attention weights by each input vector and add to get the final context vector
- Using attention, each output is able to take into consideration each input word



Attention in RNNs



We can visualize the attention weights for each word

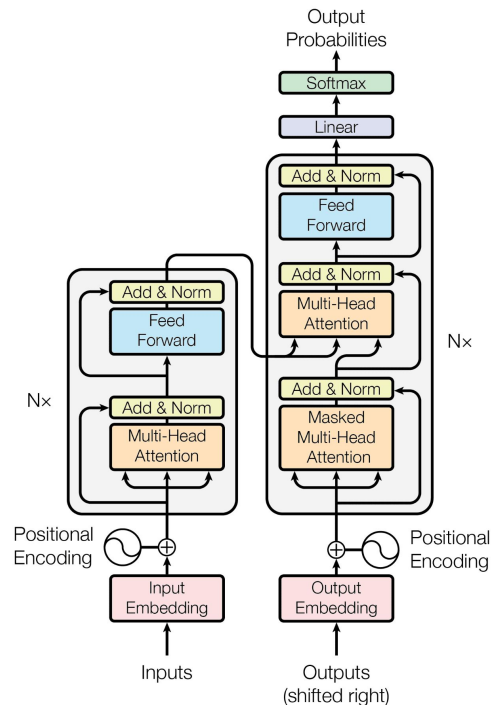


A good resource/tutorial for attention can be found here:

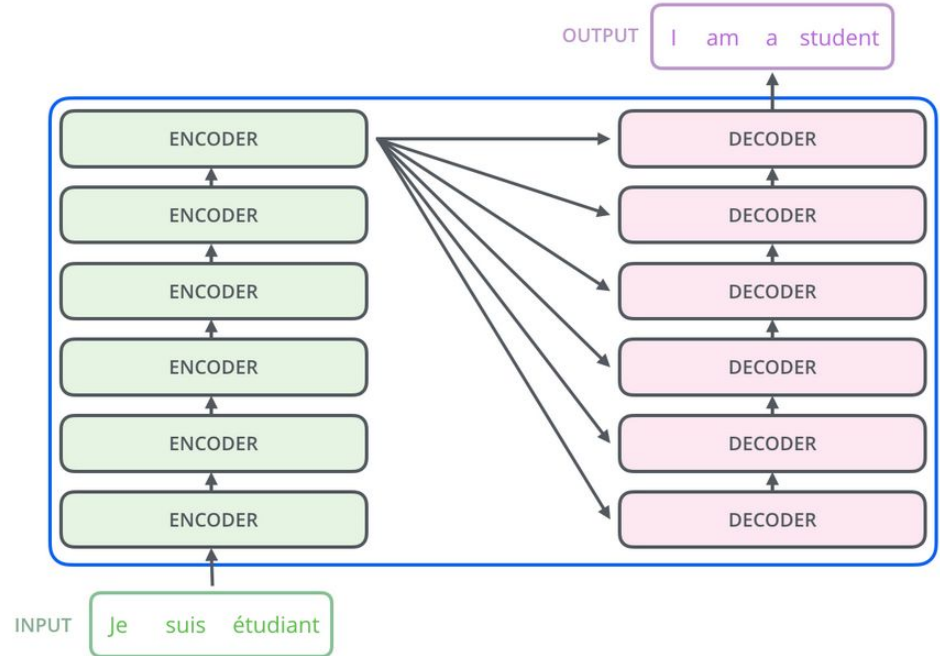

<https://jalammar.github.io/visualizing-neural-machine-translation-mechanics-of-seq2seq-models-with-attention/>

Transformers

- Transformers are also language models, but they don't use any recurrent structure unlike RNNs or LSTMs
- Rather, they introduce the idea of **self-attention**
- “Attention is All You Need” [1]



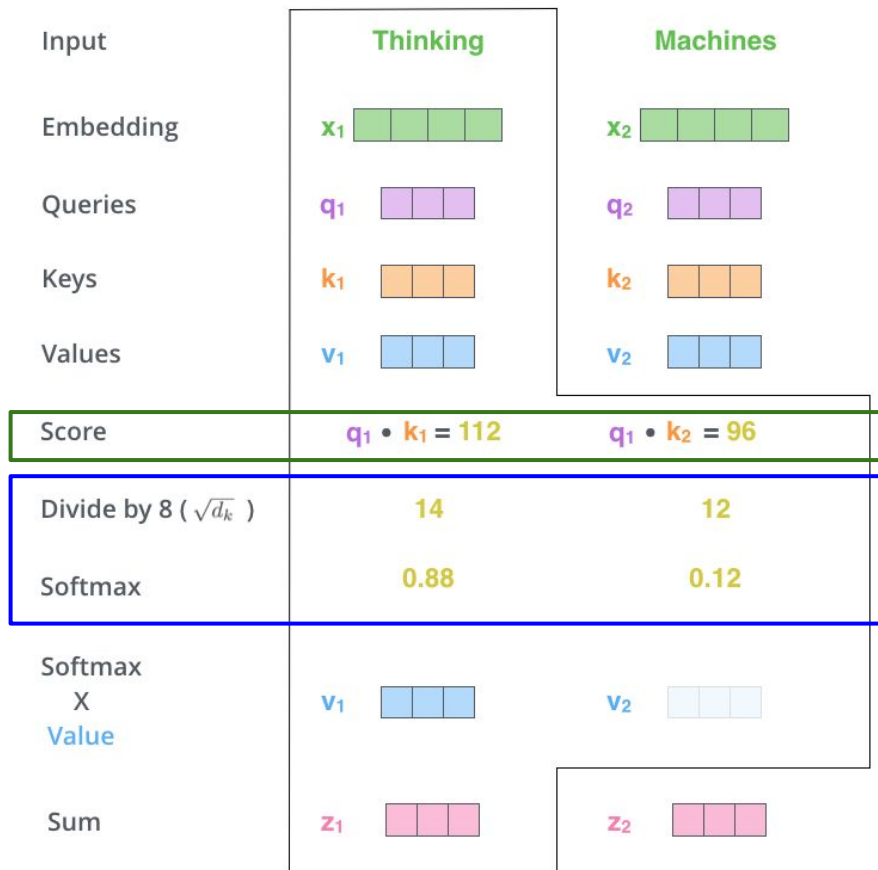
Encoder + Decoder Architecture



Similar to RNNs, transformers follow an encoder-decoder structure with stacked layers/blocks.

Self-Attention

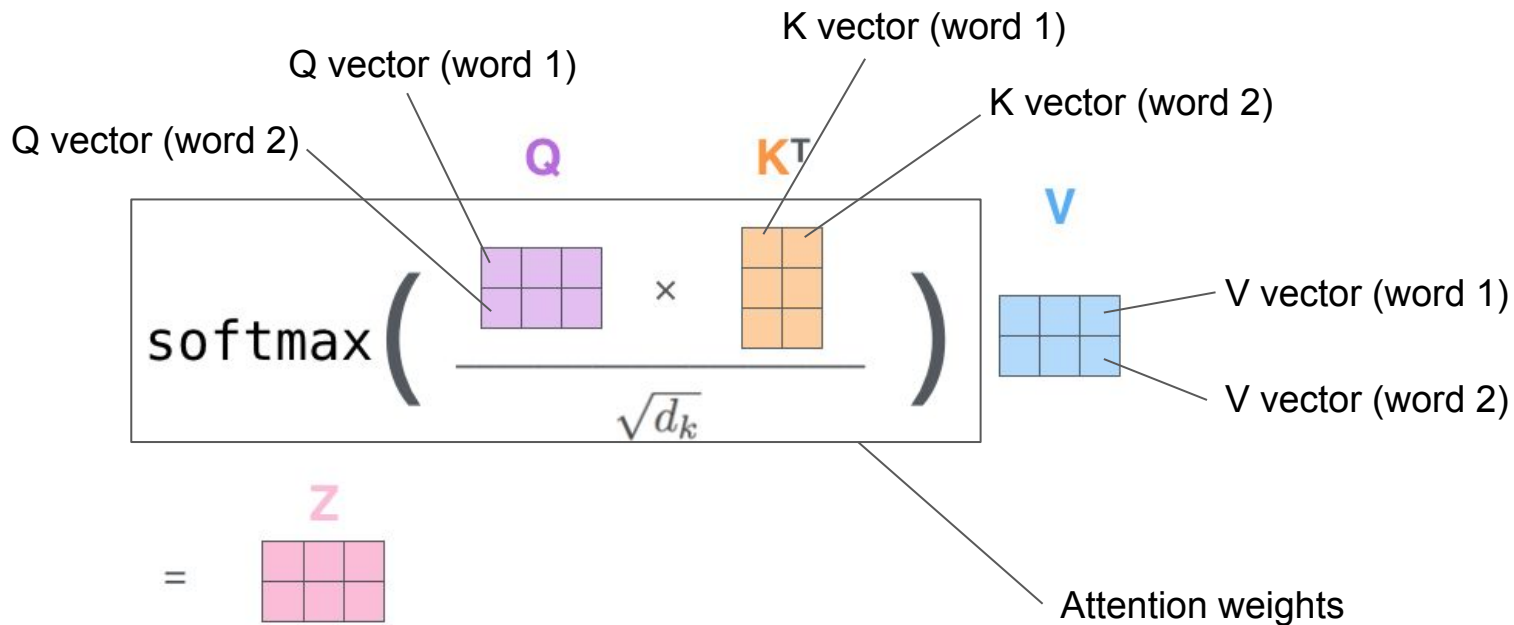
- Each input pays attention to each other input (hence “self”-attention)
- In addition to the input vector (embedding), each input also has a corresponding query, key, and value vector
- How to get the attention weights:
 - Similar to RNN – take a dot product
 - Dot product of query vector of current word with key vector of each of the other input words
 - Scale and softmax to get final weights
- How to get final output:
 - Weighted sum of value vectors (using attention weights)



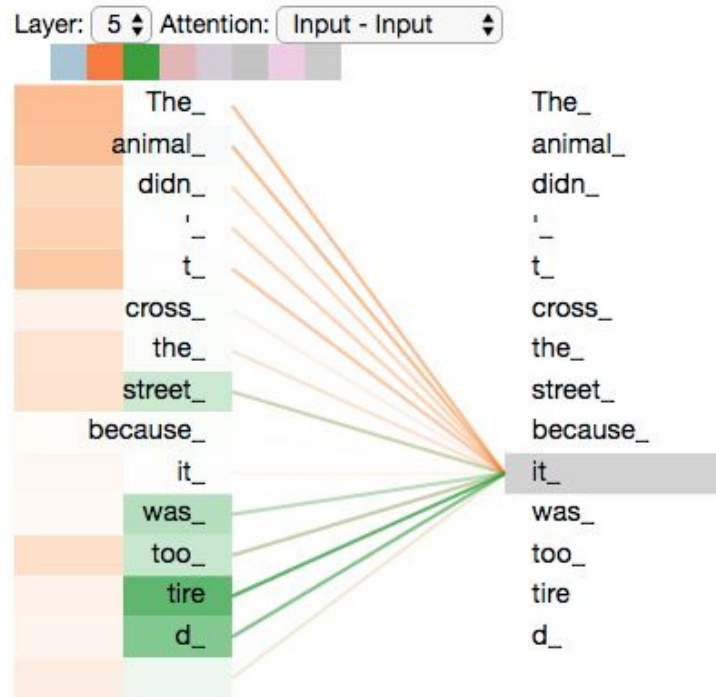
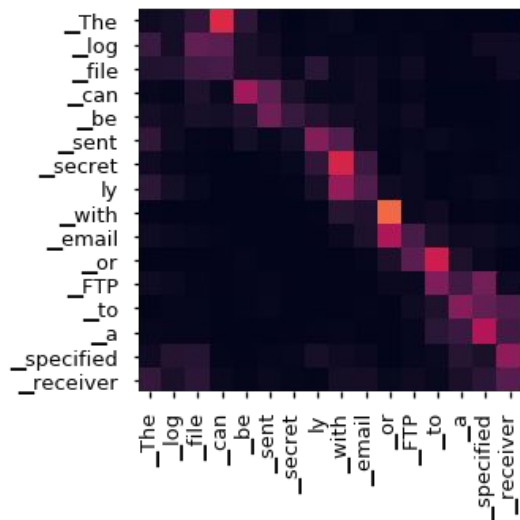
Self-Attention

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

In matrix form:



Self-Attention




We can plot the (self-)attention weights similar to how we previously plotted attention weights for RNNs

Self-Attention

- In practice, rather than directly calling $\text{Attention}(Q, K, V)$:

We instead call $\text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$ where the W s are parameter (weight) matrices (basically a linear layer)

- This allows us to have multiple “heads” for attention



Multi-Headed Attention

Z_i is the result of a single attention head i


$$\text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$

Then we concatenate Z_1, Z_2, \dots, Z_k and multiply by a final weight vector to get the final Z



1) This is our input sentence*

Thinking
Machines


2) We embed each word*





3) Split into 8 heads. We multiply X or R with weight matrices



...



4) Calculate attention using the resulting $Q/K/V$ matrices



...



5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer




...



W^O




Z



* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one



Multi-Headed Attention




Different attention heads can pay attention to different parts of the input

Decoding Masking

In the decoder, the self-attention layer is only allowed to attend to earlier positions in the output sequence. This is done by masking future positions

In practice, a “mask” is a binary matrix (0s and 1s), where the 0s represent the positions to “mask out”



Positional Encoding

- As the name suggests, we need a way to encode the position of the tokens
 - Introduce a “position matrix” to add to the input matrix
 - This “position matrix” is determined as follows:


$$\mathbf{P}_{(\text{pos}, 2i)} = \sin\left(\frac{\text{pos}}{10000^{2i/d}}\right) \quad \mathbf{P}_{(\text{pos}, 2i+1)} = \cos\left(\frac{\text{pos}}{10000^{2i/d}}\right)$$

- The final input is the sum of the position matrix and the input matrix



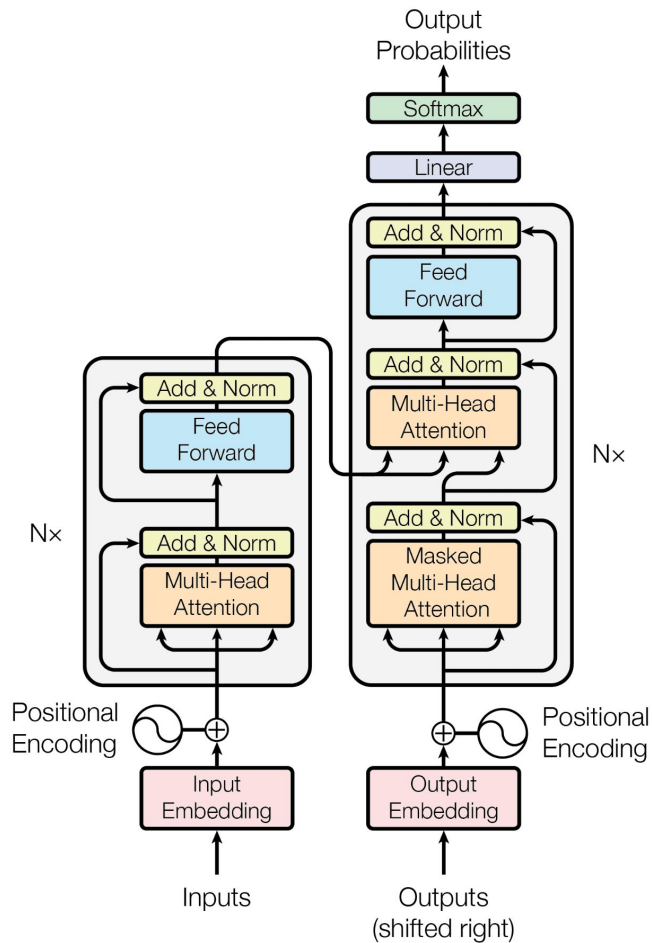
Positional Encoding

$$\mathbf{P}_{(\text{pos}, 2i)} = \sin\left(\frac{\text{pos}}{10000^{2i/d}}\right) \quad \mathbf{P}_{(\text{pos}, 2i+1)} = \cos\left(\frac{\text{pos}}{10000^{2i/d}}\right)$$



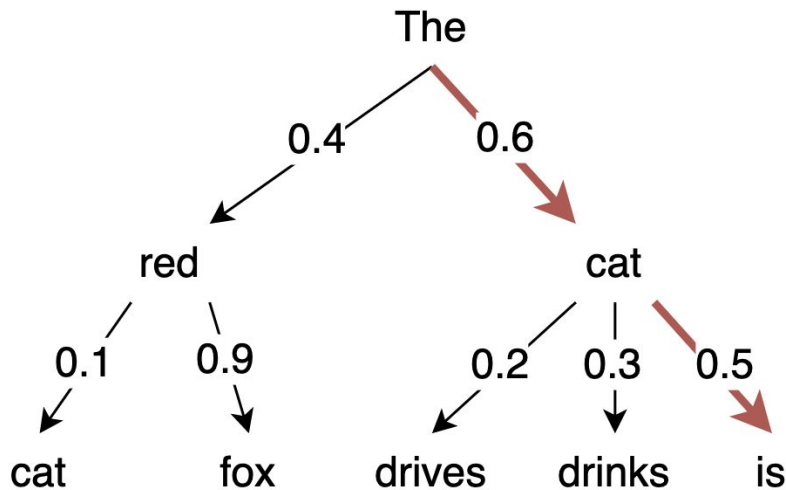
Here, pos is the index of the word in the sentence, d is the number of dimensions (i.e. embedding size), and $2i$ (or $2i+1$) is the index along the d axis.

Putting it all together



Decoding: Greedy

- Once the model is trained, how do we generate results?
 - Recall that a language model is a **probability distribution over words**.
 - We can simply take argmax at each step!



Downside of greedy decoding:


This may not necessarily lead to the most probable sequence. In this example, “The red fox” is actually more probable (0.36) than “The cat is” (0.30).

Decoding: Greedy

How do we capture the sentence with the highest probability?

We will need to search the entire tree. This can get computationally expensive (intractable) since each node will introduce V new children (where V =vocabulary size).

Can we introduce some kind of compromise?

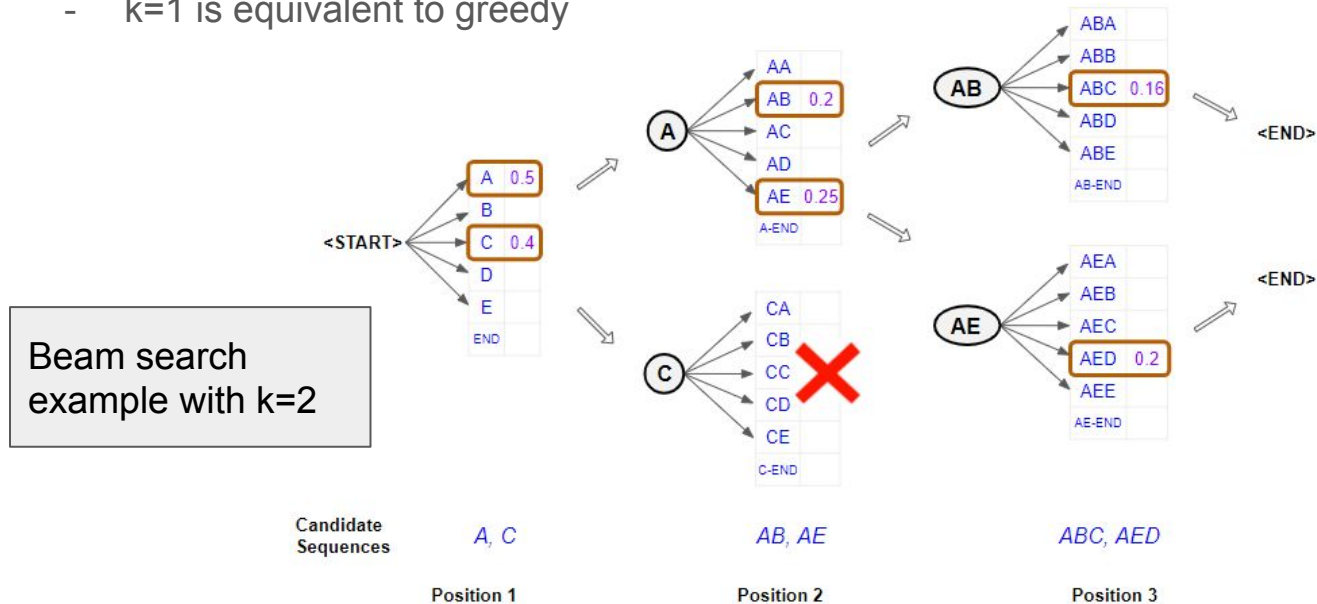


Downside of greedy decoding:

This may not necessarily lead to the most probable sequence. In this example, “The red fox” is actually more probable (0.36) than “The cat is” (0.30).

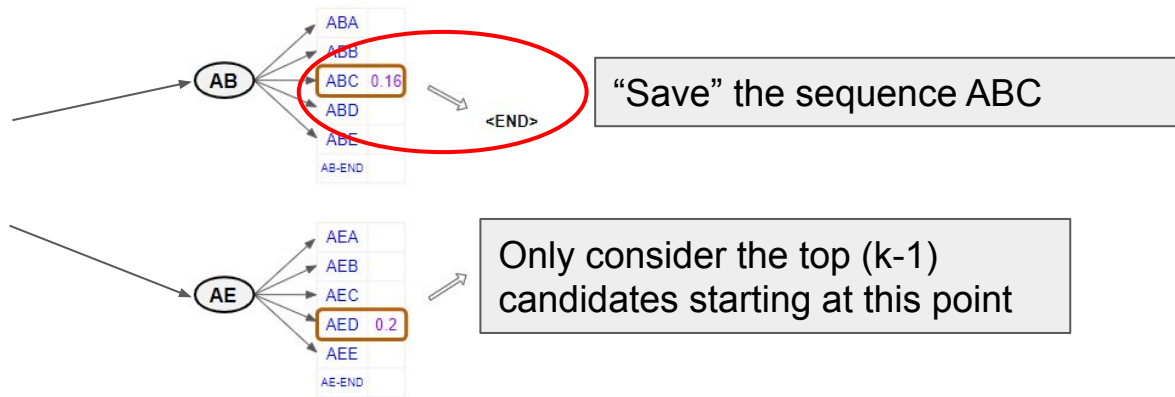
Decoding: Beam Search

- Instead of keeping only the top candidate at each step, we keep the top k candidates at each step
 - k is called the “beam width” or “beam size”
 - k=1 is equivalent to greedy



Decoding: Beam Search

- Instead of keeping only the top candidate at each step, we keep the top k candidates at each step
 - k is called the “beam width” or “beam size”
 - $k=1$ is equivalent to greedy
- What happens when we reach the end-of-sentence [EOS] token?
 - Save that sequence as a possible “final candidate”.
 - Reduce your beam size by one – only consider the top $(k-1)$ candidates starting from now



Decoding: Beam Search

- Instead of keeping only the top candidate at each step, we keep the top k candidates at each step
 - k is called the “beam width” or “beam size”
 - $k=1$ is equivalent to greedy
- What happens when we reach the end-of-sentence [EOS] token?
 - Save that sequence as a possible “final candidate”.
 - Reduce your beam size by one – only consider the top $(k-1)$ candidates starting from now
- How do we select the final answer?
 - Consider all beams from the final timestep AND all “final candidates” from the [EOS] step
 - Select the sentence with the highest probability
 - In practice, we add log-likelihood probabilities instead of multiplying probabilities (i.e. compute $\log(p_1)+\log(p_2)+\log(p_3)$ instead of $p_1p_2p_3$.)

Evaluation (BLEU Score)

- This is used to evaluate the quality of our translations
- Score between 0 and 1
- Based on **n-gram matching**
- An n-gram is a contiguous sequence of words of size n
- e.g. “The dog is a happy dog”
 - 1-grams: { [The], [dog], [is], [a], [happy], [dog] }
 - 2-grams: { [The, dog], [dog, is], [is, a], [a, happy], [happy, dog] }
 - 3-grams: { [The, dog, is], [dog, is, a], [is, a, happy], [a, happy, dog] }

BLEU Score Calculation

- We will build up towards this formula:

$$P_n(y, \hat{y}) \triangleq \frac{\sum_{x \in \{\text{unique n-grams in } \hat{y}\}} \min(C(\hat{y}, x), C(y, x))}{|\hat{y}| - n + 1}$$

BLEU Score Calculation

- Let's consider 1-grams for now.
- Target sentence: $y = \{ [\text{The}], [\text{dog}], [\text{is}], [\text{a}], [\text{happy}], [\text{dog}] \}$
- Predicted sentence: $y_{\text{hat}} = \{ [\text{The}], [\text{cat}], [\text{is}], [\text{a}], [\text{very}], [\text{happy}], [\text{cat}] \}$
- **Initial attempt (wrong):** Calculate number of n-gram overlaps between prediction and target
 - "The" = 1
 - "cat" = 0
 - "is" = 1
 - "a" = 1
 - "very" = 0
 - "happy" = 1
 - "cat" = 0
- Total = $\text{sum}(\text{counts}) / \text{total_ngrams_in_}y_{\text{hat}} = 4/7$

Why is this not a good idea?

BLEU Score Calculation

- Let's consider 1-grams for now.
- Target sentence: $y = \{ [\text{The}], [\text{dog}], [\text{is}], [\text{a}], [\text{happy}], [\text{dog}] \}$
- Predicted sentence: $y_hat = \{ [\text{happy}], [\text{happy}], [\text{happy}], [\text{happy}], [\text{happy}] \}$
- **Initial attempt (wrong)**: Calculate number of n-gram overlaps between prediction and target
 - "happy" = 1
 - "happy" = 1
 - "happy" = 1
 - "happy" = 1
 - "happy" = 1
- Total = $\text{sum}(\text{counts}) / \text{total_ngrams_in_y_hat} = 5/5$

We need to find a way to deal with repeated words.

BLEU Score: Clipped Precision

- Previously:

For **each n-gram** in the predicted sentence (y_{hat}), count the **number of times it appears in the target sentence (y)**. Take the sum of these counts and divide by $\text{total_ngrams_in_}y_{\text{hat}}$

- With clipped precision:

For **each unique n-gram** in the predicted sentence (y_{hat}), count the **$\min(\text{number of times it appears in } y, \text{number of times it appears in } y_{\text{hat}})$** . Take the sum of these counts and divide by $\text{total_ngrams_in_}y_{\text{hat}}$

BLEU Score: Clipped Precision

$$P_n(y, \hat{y}) \triangleq \frac{\sum_{x \in \{\text{unique n-grams in } \hat{y}\}} \min(C(\hat{y}, x), C(y, x))}{|\hat{y}| - n + 1}$$

- With clipped precision:

For **each unique n-gram** in the predicted sentence (\hat{y}), count the **min(number of times it appears in y , number of times it appears in \hat{y})**. Take the sum of these counts and **divide by total_ngrams_in_y_hat**

BLEU Score Calculation

- Let's consider 1-grams for now.
- Target sentence: $y = \{ [\text{The}], [\text{dog}], [\text{is}], [\text{a}], [\text{happy}], [\text{dog}] \}$
- Predicted sentence: $y_{\text{hat}} = \{ [\text{happy}], [\text{happy}], [\text{happy}], [\text{happy}], [\text{happy}] \}$
- **With clipped precision:**
 - "happy" = $\min(1, 5) = 1$
- Total = $\text{sum}(\text{counts}) / \text{total_ngrams_in_}y_{\text{hat}} = 1/5$

Are we done? Is this good enough?

BLEU Score Calculation

- Let's consider 1-grams for now.
- Target sentence: $y = \{ [\text{The}], [\text{dog}], [\text{is}], [\text{a}], [\text{happy}], [\text{dog}] \}$
- Predicted sentence: $y_{\text{hat}} = \{ [\text{happy}], [\text{happy}], [\text{happy}], [\text{happy}], [\text{happy}] \}$
- **With clipped precision:**
 - "happy" = $\min(1, 5) = 1$
- Total = $\text{sum}(\text{counts}) / \text{total_ngrams_in_}y_{\text{hat}} = 1/5$

- Predicted sentence: $y_{\text{hat}} = \{ [\text{happy}] \}$
- **With clipped precision:**
 - "happy" = $\min(1, 1) = 1$
- Total = $\text{sum}(\text{counts}) / \text{total_ngrams_in_}y_{\text{hat}} = 1/1$

How do we deal with short predictions?

BLEU Score: Brevity Penalty

- We introduce a **brevity penalty**
- If prediction is shorter than the target:

$$\text{penalty} = \exp(1 - (\text{len_target} / \text{len_prediction}))$$

- Otherwise, no penalty:

$$\text{penalty} = 1$$

- We multiply this penalty with our BLEU score

BLEU Score Calculation

- Let's consider 1-grams for now.
- Target sentence: $y = \{ [\text{The}], [\text{dog}], [\text{is}], [\text{a}], [\text{happy}], [\text{dog}] \}$
- Predicted sentence: $y_{\text{hat}} = \{ [\text{happy}] \}$
- **With clipped precision and brevity penalty:**
 - "happy" = $\min(1, 1) = 1$
- Total = $[\text{sum}(\text{counts}) / \text{total_ngrams_in_}y_{\text{hat}}] * \text{brevity_penalty}$
 $= (1/1) * \exp(1 - 6/1) = 1 * \exp(-5) = \mathbf{0.0067}$

BLEU Score: Final Calculation

- Previously, we were only considering 1-grams.
- Repeat this same process for 2-grams, 3-grams, ... until k-grams.
- This gives scores $P_1, P_2, \dots P_k$
- The final BLEU score is the **geometric mean** of $(P_1, P_2, \dots P_k)$

BLEU Score: Worked Example

- Target sentence: {"the", "fat", "cat", "ate", "the", "fat", "rat"}
- Predicted sentence: {"the", "fat", "cat", "ate", "the", "cat"}
- Suppose $k=3$. We use "c()" here to denote the clipped precision
- **P1**: $c(["the"]) = \min(2, 2)$, $c(["fat"]) = \min(2, 1)$, $c(["cat"]) = \min(1, 2)$, $c(["ate"]) = \min(1, 1)$
- **P1** = $(2+1+1+1) / \text{total_num_1_grams_in_pred} = \mathbf{5/6}$
- **P2**: $c(["the\ fat"]) = 1$, $c(["fat\ cat"]) = 1$, $c(["cat\ ate"]) = 1$, $c(["ate\ the"]) = 1$, $c(["the\ cat"]) = 0$
- **P2** = $(1+1+1+1+0) / \text{total_num_2_grams_in_pred} = \mathbf{4/5}$
- **P3**: $c(["the\ fat\ cat"]) = 1$, $c(["fat\ cat\ ate"]) = 1$, $c(["cat\ ate\ the"]) = 1$, $c(["ate\ the\ cat"]) = 0$
- **P3** = $(1+1+1+0) / \text{total_num_3_grams_in_pred} = \mathbf{3/4}$
- Final score = $\text{geometric_mean}(\mathbf{5/6}, \mathbf{4/5}, \mathbf{3/4}) * \text{brevity_penalty}$
 $= (3/6)^{(1/3)} * \exp(1 - 7/6) = \mathbf{0.672}$

Programming Portion: Auto-grader

1. Positional Encoding
2. Self-attention Layer
3. Lookahead Mask
4. Beam Search Prediction
5. BLEU Score (617 only)

Programming Portion: Experiments

1. (autograder)
2. Plotting train and test loss
 - This should be relatively straightforward (similar to previous assignments)
 - Provided function: `train()`
 - Make sure to save your models, as you will be loading/using them in the next questions
3. Decoding (beam search)
 - Task: Generate translations for a few sentences and report the generations
 - Provided function: `decode_sentence()`
4. Visualizing Attention
 - Provided function: `visualize_attention()`
 - How to generate attention matrix: output of `Transformer.forward`

Programming Portion: Experiments

5. Beam Search

- Task: Investigate the effect of different beam sizes

6. BLEU Score (617 only)

- Generate output first
- Then take `BLEU_Score(output_prediction, ground_truth_target)`